

# Question 1

Jin Barai


```
library(carData)
data(TitanicSurvival)
Titanic = na.omit(TitanicSurvival)
Titanic = Titanic[Titanic$sex == "female", ]
Titanic$survived1 = as.numeric(Titanic$survived == "yes")
head(Titanic)
```

##	survived	sex	age	passengerClass	survived1
## Allen, Miss. Elisabeth Walton	yes	female	29	1st	1
## Allison, Miss. Helen Loraine	no	female	2	1st	0
## Allison, Mrs. Hudson J C (Bessi	no	female	25	1st	0
## Andrews, Miss. Kornelia Theodos	yes	female	63	1st	1
## Appleton, Mrs. Edward Dale (Cha	yes	female	53	1st	1
## Astor, Mrs. John Jacob (Madelei	yes	female	18	1st	1


(a)

```
par(mfcol = c(1,2))
age = Titanic$age
rx = range(age)
hist(age, breaks=seq(0, 80, 8), prob=TRUE, xlab="Female passenger's Age", main="Bin width equal to 8 ye
hist(age, breaks=quantile(age, p=seq(0,1,length.out=11)), prob=TRUE, main="Varying bin width", xlab="Fer
```

**Bin width equal to 8 years**



**Varying bin width**



(b)



```
survived = Titanic$survived1
```

```
plot.fun <- function (agegroup){
  plot(age, survived, pch=19,col=c(adjustcolor("firebrick",0.5), adjustcolor("blue", 0.5))[survived +1],
    xlim=c(0,80), xlab="Age", ylab="Survival")
  x = matrix(0, nrow=length(agegroup)-1, ncol=5)
  dimnames(x)[[2]] = c("Lower", "Upper", "Total", "Num.Survived", "Prop.Survived")
  for (i in 1:nrow(x)) {
    x[i,1:2] = c(agegroup[i], agegroup[i+1])
    x[i,3] = sum(age > agegroup[i] & age <= agegroup[i+1])
    x[i,4] = sum(survived[age > agegroup[i] & age <= agegroup[i+1]] )
    x[i,5] = round(x[i,4]/x[i,3],3)
  }
  points(agegroup[-length(agegroup)]+ diff(agegroup)/2, x[,5], pch=19, col=adjustcolor("black", 0.5),type="n",
    xlab="Age", ylab="Survival")
}
```

```
par(mfcol = c(1,2))
agegroup1 = seq(0,80,8)
agegroup2 = quantile(age, seq(0,1,length.out=11))
agegroup2[1] = 0
plot.fun(agegroup1)
```

##		Lower	Upper	Total	Num.Survived	Prop.Survived
##	[1,]	0	8	34	24	0.706
##	[2,]	8	16	30	20	0.667
##	[3,]	16	24	108	80	0.741
##	[4,]	24	32	77	52	0.675
##	[5,]	32	40	61	51	0.836
##	[6,]	40	48	38	29	0.763
##	[7,]	48	56	23	22	0.957
##	[8,]	56	64	16	13	0.812
##	[9,]	64	72	0	0	NaN
##	[10,]	72	80	1	1	1.000

```
plot.fun(agegroup2)
```




##		Lower	Upper	Total	Num.Survived	Prop.Survived
##	[1,]	0.0	9.7	39	25	0.641
##	[2,]	9.7	18.0	53	38	0.717
##	[3,]	18.0	21.0	27	18	0.667
##	[4,]	21.0	24.0	53	43	0.811
##	[5,]	24.0	27.0	25	16	0.640
##	[6,]	27.0	30.0	38	26	0.684
##	[7,]	30.0	35.0	39	34	0.872
##	[8,]	35.0	40.6	36	27	0.750
##	[9,]	40.6	49.0	40	31	0.775
##	[10,]	49.0	76.0	38	34	0.895

We would prefer the age range with varying lengths. This is because there are roughly the same number of units within each age group. For Equal bin widths, the total number of units in each age group varies significantly.

(c)

i.

```
probit.fn <- function(z) {  
  pnorm(z)  
}  
  
z = seq(-6,6,0.1)  
  
plot(z,probit.fn(z),type='l')
```




ii.

```
plot.fun(agegroup2)
```

##		Lower	Upper	Total	Num.Survived	Prop.Survived
##	[1,]	0.0	9.7	39	25	0.641
##	[2,]	9.7	18.0	53	38	0.717
##	[3,]	18.0	21.0	27	18	0.667
##	[4,]	21.0	24.0	53	43	0.811

```
## [5,] 24.0 27.0 25 16 0.640
## [6,] 27.0 30.0 38 26 0.684
## [7,] 30.0 35.0 39 34 0.872
## [8,] 35.0 40.6 36 27 0.750
## [9,] 40.6 49.0 40 31 0.775
## [10,] 49.0 76.0 38 34 0.895
```

```
z = seq(0,300,0.1)
lines(z,probit.fn(-1/2 + 0.03*z))
```



(d)

The log-likelihood is given by

$$l(\theta) = l(\alpha, \beta) = \sum_{i=1}^N \left[ y_i \log \frac{p_i}{1 - p_i} + \log(1 - p_i) \right]$$

where

$$p_i = \Phi(\hat{y}) = \Phi(\alpha + \beta[x_i - \bar{x}]) = \frac{1}{\sqrt{2\pi}} e^{-\frac{(\alpha + \beta[x_i - \bar{x}])^2}{2}}$$

Now we proceed by finding partial derivatives with respect to  $\alpha$  and  $\beta$

$$\frac{\partial l}{\partial \alpha} = \sum_{i=1}^N \frac{\partial l}{\partial p_i} \frac{\partial p_i}{\partial z} \frac{\partial z}{\partial \alpha}$$

$$\frac{\partial l}{\partial \beta} = \sum_{i=1}^N \frac{\partial l}{\partial p_i} \frac{\partial p_i}{\partial z} \frac{\partial z}{\partial \beta}$$

$$\frac{\partial l}{\partial p_i} = \frac{\partial}{\partial p_i} \left[ y_i \log \frac{p_i}{1-p_i} - \frac{1}{1-p_i} \right] = \frac{y_i - p_i}{p_i(1-p_i)}$$

$$\Phi(z) = \int_{-\infty}^z \frac{1}{\sqrt{2\pi}} e^{-u^2/2} \Rightarrow \Phi'(z) = \phi(z) = \frac{1}{\sqrt{2\pi}} e^{-z^2/2}$$

$$p = \int_{-\infty}^z \frac{1}{\sqrt{2\pi}} e^{-u^2/2} \Rightarrow \frac{\partial p}{\partial z} = \frac{1}{\sqrt{2\pi}} e^{-z^2/2} = \phi(z)$$

$$\frac{\partial z}{\partial \alpha} = \frac{\partial}{\partial \beta} (\alpha + \beta[x_i - \bar{x}]) = 1$$

$$\frac{\partial z}{\partial \beta} = \frac{\partial}{\partial \beta} (\alpha + \beta[x_i - \bar{x}]) = (x_i - \bar{x})$$

$$\frac{\partial l}{\partial \alpha} = \sum_{i=1}^N \frac{\partial l}{\partial p_i} \frac{\partial p_i}{\partial z} \frac{\partial z}{\partial \alpha} = \sum_{i=1}^N \frac{y - p_i}{p_i(1-p_i)} \times \phi(\hat{y}) \times 1$$

$$\frac{\partial l}{\partial \beta} = \sum_{i=1}^N \frac{\partial l}{\partial p_i} \frac{\partial p_i}{\partial z} \frac{\partial z}{\partial \beta} = \sum_{i=1}^N \frac{y - p_i}{p_i(1-p_i)} \times \phi(\hat{y}) \times (x_i - \bar{x})$$

$$\frac{\partial l}{\partial (\alpha, \beta)} = \sum_{i=1}^N \frac{\partial l}{\partial p_i} \frac{\partial p_i}{\partial z} \frac{\partial z}{\partial \alpha} + \sum_{i=1}^N \frac{\partial l}{\partial p_i} \frac{\partial p_i}{\partial z} \frac{\partial z}{\partial \beta} = \sum_{i=1}^N \frac{y - p_i}{p_i(1-p_i)} \times \phi(\hat{y}) \times \begin{bmatrix} 1 \\ (x_i - \bar{x}) \end{bmatrix}$$

(e)

i.

```
createObjBinary <- function(x,y) {
  ## local variable
  xbar <- mean(x)
  ## Return this function
  function(theta) {
    alpha <- theta[1]
    beta <- theta[2]
    y_hat = alpha + beta * (x - xbar)
    pi = probit.fn( y_hat )
    -1*sum( y*log(pi/(1-pi)) + log(1-pi) )
  }
}
```

ii.

```
createBinaryLogisticGradient <- function(x,y) {  
  ## local variables  
  xbar <- mean(x)  
  ybar <- mean(y)  
  N <- length(x)  
  function(theta) {  
    alpha <- theta[1]  
    beta <- theta[2]  
    y_hat = alpha + beta * (x - xbar)  
    pi = probit.fn( y_hat )  
    resids = y - pi  
    -1*c( sum(resids), sum((x - xbar) * resids) )  
  }  
}
```

iii.

```
gradientDescent <- function(theta = 0,  
                             rhoFn, gradientFn,  
                             lineSearchFn, testConvergenceFn,  
                             maxIterations = 100, # maximum number of iterations  
                             # in gradient descent loop  
                             #  
                             tolerance = 1E-6, # parameters for the test  
                             relative = FALSE, # for convergence function  
                             #  
                             lambdaStepsize = 0.01, # parameters for the line search  
                             lambdaMax = 0.5 # to determine lambda  
                             ) {  
  
  ## Initialize  
  converged <- FALSE  
  i <- 0  
  ## LOOP  
  while (!converged & i <= maxIterations) {  
    ## gradient  
    g <- gradientFn(theta)  
    ## gradient direction  
    glength <- sqrt(sum(g^2))  
    if (glength > 0) g <- g /glength  
  
    ## line search for lambda  
    lambda <- lineSearchFn(theta, rhoFn, g,  
                           lambdaStepsize = lambdaStepsize,  
                           lambdaMax = lambdaMax)  
  
    ## Update theta  
    thetaNew <- theta - lambda * g  
    ##  
    ## Check convergence  
    converged <- testConvergenceFn(thetaNew, theta,  
                                   tolerance = tolerance,  
                                   relative = relative)  
    ## Update
```

```

    theta <- thetaNew
    i <- i + 1
  }
  ## Return last value and whether converged or not
  list(theta = theta,
        converged = converged,
        iteration = i,
        fnValue = rhoFn(theta)
      )
}

### line searching could be done as a simple grid search
gridLineSearch <- function(theta, rhoFn, g,
                           lambdaStepsize = 0.01,
                           lambdaMax = 1) {
  ## grid of lambda values to search
  lambdas <- seq(from = 0,
                 by = lambdaStepsize,
                 to = lambdaMax)

  ## line search
  rhoVals <- Map(function(lambda) {rhoFn(theta - lambda * g)},
                 lambdas)
  ## Return the lambda that gave the minimum
  lambdas[which.min(rhoVals)]
}

### Where testConvergence might be (relative or absolute)
testConvergence <- function(thetaNew, thetaOld, tolerance = 1E-10, relative=FALSE) {
  sum(abs(thetaNew - thetaOld)) < if (relative) tolerance * sum(abs(thetaOld)) else tolerance
}

gradient <- createBinaryLogisticGradient(age, survived)
rho <- createObjBinary(age, survived )

# c(-1.32,0.02609972 )
result <- gradientDescent(theta = c(0, 0),
                          rhoFn = rho, gradientFn = gradient,
                          lineSearchFn = gridLineSearch,
                          testConvergenceFn = testConvergence,
                          lambdaStepsize = 0.0001,
                          lambdaMax = 0.01,
                          maxIterations = 10^5)

### Print the results
Map(function(x){if (is.numeric(x)) round(x,4) else x}, result)

## $theta
## [1] 0.6787 0.0131
##
## $converged
## [1] TRUE
##
## $iteration
## [1] 161

```



```
##
## $fnValue
## [1] 213.4404
```

The algorithm converges.

iv.

- If age has no effect on survival then  $\beta = 0$
- If the proportion of females who survived the titanic is 0.717 then

$$0.717162 = \Phi(\alpha) = \int_{-\infty}^{\alpha} \frac{1}{\sqrt{2\pi}} e^{-u^2/2} \Rightarrow \alpha = 0.574$$

Hence the starting value is  $(\alpha, \beta) = (0.574, 0)$

```
result <- gradientDescent(theta = c(0.574, 0),
                           rhoFn = rho, gradientFn = gradient,
                           lineSearchFn = gridLineSearch,
                           testConvergenceFn = testConvergence,
                           lambdaStepsize = 0.0001,
                           lambdaMax = 0.01,
                           maxIterations = 10^5)

### Print the results
Map(function(x){if (is.numeric(x)) round(x,4) else x}, result)
```

```
## $theta
## [1] 0.6898 0.0133
##
## $converged
## [1] TRUE
##
## $iteration
## [1] 76
##
## $fnValue
## [1] 213.4148
```

The number of iterations went from 161 to 76, hence an improvement.

(f)

i.

```
temp = plot.fun(agegroup2)
z = seq(0, 80, length.out=100)
lines(z, probit.fn(0.6898 + 0.0133*(z-mean(age))))
```

ii.

```
x = apply(temp[,1:2],1,mean)
probit.probit.prop = probit.fn(0.6898 + 0.0133*(x - mean(age)))
propx = cbind(temp, probit.probit.prop)
round(propx,3)
```

iii.

- The parametric model assumes monotonic increasing or decreasing survival proportion
- The non-parametric model assumes constant survival proportions over each interval

iv.

$$p = \int_{-\infty}^z \frac{1}{\sqrt{2\pi}} e^{-u^2/2}$$

where  $z = 0.6898 + 0.0133 \times (x - \bar{x})$ . When there is a 50-50 chance  $p = 1/2$  so we have

$$1/2 = \int_{-\infty}^z \frac{1}{\sqrt{2\pi}} e^{-u^2/2} \Rightarrow z = 0$$

$$0 = 0.6898 + 0.0133 \times (x - \bar{x}) \Rightarrow x = \bar{x} - \frac{0.6898}{0.0133} = 28.68707 - \frac{0.6898}{0.0133}$$